

FORMA

You have the design in your head.
Forma writes it.

A **Grammarly-style overlay** for vibe coders. As you type vague UI language on v0, Cursor, and Lovable, Forma rewrites it into canonical components with animated previews — so the AI builder generates **what you actually wanted**.

Powered by an asymmetric inference architecture only AMD MI300X enables.

RYAN ZHANG · CHRISTIAN LEE

Llama 3.1 8B + 70B AWQ-INT4 · vLLM · ROCm 7.0 · 7-agent consensus on one MI300X

AI builders got better. Prompts didn't.

v0, Cursor, Lovable, Bolt, base44 — all compete on output quality. **None of them fix the input.** And the input is broken at a level no model can solve: **vocabulary.**

▼ THE CAUSE

UI patterns are visual. Prompts are verbal.

60

CANONICAL UI PATTERNS
FORMA RECOGNIZES

<1s

TO SURFACE THE RIGHT ONE
AS YOU TYPE

Designers and PMs see the design in their head — but the right words ("Glassmorphic Popover", "Off-Canvas Drawer") are specialized vocabulary they were never taught. **You can't type a word you don't know.**

▼ THE CONSEQUENCE


You burn three regenerations getting v0 to make what was already in your head.

The model isn't wrong — it's working from the words you typed, not the design you meant. So you iterate. Tweak. Regenerate. Tweak again. Designers and PMs lose hours to a vocabulary problem they don't know they have.

Every regeneration is a vocabulary tax.

Same builder. Same model. Forma is the difference.

Same user. Same v0.app. Same generation model. Same project. The only difference is whether they typed it — or Forma rewrote it.




SCORE

60

↓

95

+35 LIFT



Same prompt intent. Same v0.app. Different prompt quality. Different result.

You don't need the right word. You need to point at it.


UI patterns are visual. English is verbal. As you type, the 8B model finds vague phrases in under a second — then hovering reveals the canonical alternates.

1

▼ TYPE · 8B WATCHING EVERY KEYSTROKE

Score appears in under a second.

User types into vO.app. Forma reads the textarea, scores the prompt 0–100, and surfaces a **Run Full Analysis Pro** action.




2

▼ HOVER · THE MAGIC MOMENT

Tooltip with main term + 3 alternates.

Each alternate has its own animated preview. Click the one that matches the picture in your head — vocabulary is no longer a barrier.



Pro mode adds motion specs and accessibility.


Toggle Pro and Forma expands the canonical term into a full prompt: width, easing curve, focus management, ARIA attributes. One click, score 60 → 95.

3

▼ PRO MODE · FULL CANONICAL EXPANSION

From noun to spec.

"Off-Canvas Drawer" becomes "slides from right edge, 320px width, 250ms ease-in-out, semi-transparent backdrop" — the exact prompt v0.app needs.




4

▼ ACCEPT · SCORE JUMPS GREEN

60 → 95 in one click.

Textarea is rewritten with the full Glassmorphic Popover spec. Score turns green · v0.app generates what you actually wanted.



8.75× parameter asymmetry. One GPU. Two business models.

Two production-grade Llamas — one ten times the size of the other — sharing the same MI300X. Free tier on the 8B; paid tier on the 70B. The hardware choice and the freemium business model are the same decision.

FREE TIER · PORT 30000

Llama 3.1 8B Instruct

Detection + scoring on every keystroke · powers the inline tooltip

~800ms
LATENCY

15GiB
VRAM

44×
CONCURRENCY

PRO TIER · PORT 8000

Llama 3.1 70B AWQ-INT4

7-agent consensus on Run Full Analysis · the moat for paying users

~40s
DEEP LATENCY

37GiB
VRAM

21×
CONCURRENCY



On H100, freemium doesn't pencil.

This isn't an architecture preference. It's the only configuration where the free-to-paid economics work. H100 forces two GPUs. Two GPUs kill freemium.

NVIDIA H100

80 GiB HBM3 • 3.35 TB/s

VRAM ALLOCATION ATTEMPT — SINGLE GPU

8B • 15

70B • 37

KV STARVED

GPU INSTANCES NEEDED

2×

COST PER FREE REQUEST

~2× higher

Both models technically fit. ~28 GiB headroom is not enough KV cache for usable concurrency. Forces a two-GPU topology — one model per GPU — doubling infra cost and adding cross-device latency to every paid request.

AMD MI300X

192 GiB HBM3 • 5.3 TB/s

ACTUAL DEPLOYED ALLOCATION — PRODUCTION

8B

70B

~140 GiB • KV + PARALLEL

GPU INSTANCES NEEDED

1×

COST PER FREE REQUEST

~\$0.01 /1K reqs


Both tiers in same VRAM. 140 GiB headroom absorbs 11.2 GiB realtime KV + 26.25 GiB deep KV plus parallel users. Freemium isn't a marketing decision — it's the architecture this hardware uniquely enables.

Cost derivation: \$1.99/hr droplet ÷ 3600s ÷ (44 concurrent × ~800ms per req) ≈ \$0.00001/req for the 8B free tier

07 / 11

From typed character to scored prompt in 800ms.

Five layers, two networks, one MI300X. Every layer is real and deployed — judges can verify each by hitting the URLs on the closing slide.



▼ REAL PRODUCTION STACK Inference: vLLM 0.17.1 Compute: ROCm 7.0 Quant: AWQ-INT4 (70B) Backend: FastAPI on Railway

GPU: MI300X • 192 GiB • 5.3 TB/s

Five layers • two networks • one MI300X • ~800ms end-to-end • this is what runs on every keystroke.

Why this can't be a single LLM call.

Each agent owns distinct context, prompt, output schema. **5 of 6 agents aligned** on the right answer in the screenshot below — that's not a single 70B inference, that's six independent votes feeding an arbitrator. Completed in 39.9s on production MI300X.

▼ THE PIPELINE · 7 AGENTS ON MI300X

 **Detector**

→ finds "popup that slides in" · 1.2s on 8B

 **Critic**

→ scores prompt · weakness: ambiguous motion

 **Reformulator**

→ primary: Glassmorphic Popover; alts: 3 candidates

 **Style**

→ user prefers modern, sleek components

 **Memory**

→ user history: tends toward Glassmorphic

 **Coach**

→ adds aria-label + focus management

 **Consensus**

→ arbitrates 6 outputs · score 90 · 5/6 aligned

▲ LIVE CONSENSUS OUTPUT · PRODUCTION MI300X

What do you want to create

from the side

Run Full Analysis PRO

Deep Analysis

7-AGENT MULTI-AGENT CONSENSUS

Llama 3.1 70B AWQ · AMD MI300X 192GB HBM3 · vLLM 0.17.1
ROCm 7.0
Completed in 39.9s

CONSENSUS

90

Design a Glassmorphic Popover with a slide-in animation from the right edge of the viewport, using Framers Motion's AnimatePresence and shaden/ui for layout management, and implement accessibility features.

5 of 6 agents aligned

The agents are in agreement on the design of a Glassmorphic Popover with a slide-in animation. The Reformulator's suggested design aligns with the user's style preferences, and the Memory agent indicates that the user tends to prefer modern, sleek components like the Glassmorphic Popover. The Coach agent's suggestions for adding accessibility features also align with the Reformulator's suggestions, increasing the confidence in the recommendation.

✓ Accept Reformulated Prompt

70B AWQ · 192 GiB HBM3 · vLLM 0.17.1 · ROCm 7.0 · completed in 39.9s

69 / 11

Translation is the wedge. Memory is the platform.

Every prompt teaches Forma. Every accept narrows the user's style. **Switching costs grow with usage** — Forma sits one layer above any single AI builder, owning the relationship between human design intent and machine output.

▲ TODAY • WORKING

LIVE IN PRODUCTION

The wedge

One Chrome extension. Five AI builders. Sub-second prompt rewrites.

v0, Cursor, Replit, Bolt, base44 — all share the same textarea problem. Forma reads, scores, rewrites.

The 8B inference path is live in production right now.

800ms

8B per keystroke

60 terms

canonical UI vocabulary

5 sites

builders supported

▲ NEXT QUARTER • V2

IN DEVELOPMENT

Real cross-builder data

Forma starts learning *your* style — not the average user's.

Cross-builder activity capture replaces the illustrative data driving Style and Memory today.

Real personalization from **your** accept history, not seeded mocks.

30 prompts

preferred patterns emerge

100 prompts

brand palette inferred

ongoing

evolving design memory

▲ NEXT YEAR • V3

ROADMAP

Platform

Forma Memory API — the design layer every builder reads from.

Builders compete on output. **Forma owns the input.**

Memory API lets v0, Cursor, Lovable, Bolt read user style profiles and ship better defaults out of the gate.

API

builders integrate at SDK level

Workspace

Notion • Linear • Figma

Enterprise

self-hosted, custom canon

Builders compete on output quality. The prompt-quality layer captures more value than any single builder below it.

Built on one MI300X over one weekend.

Working extension. Deployed backend. Real models. Live demo URL is one click away — the judge can verify everything in this deck right now.

MEASURED IMPACT · ON THE LIVE EXTENSION

60 → 95

Vague description → production-ready spec, in under a second. Forma Score (0–100) measures prompt specificity; vague phrases drop the score, canonical components raise it.

free tier · 8B · ~800ms per keystroke · canonical-60 vocabulary

TEAM



Ryan Zhang
Backend · Agents · Architecture
@ryz002-rz



Christian Lee
Frontend · Demo · Design
@haminxx

Thank you, AMD. *The MI300X didn't just enable Forma — it's the only chip our business model fits on.*

TRY IT NOW · 30 SECONDS

► INSTALL FORMA · DEMO · INSTRUCTIONS
forma-ten-bay.vercel.app

▼ AMD ARCHITECTURE DEEP-DIVE
forma-production-c800.up.railway.app/amd

LIVE BACKEND · TRY THE EXTENSION
forma-production-c800.up.railway.app

REPO · MIT LICENSE
github.com/haminxx/Forma

DEVPOST SUBMISSION
devpost.com/software/forma

▼ INSTALL FORMA

Full demo, download, and instructions at forma-ten-bay.vercel.app